Logistic Regression

Step 1: Function Set

We want to find $P_{w,b}(C_1|x)$

If $P_{w,b}(C_1|x) \geq 0.5$, output C$_1$

Otherwise, output C$_2$

$$P_{w,b}(C_1|x) = \sigma(z)$$

$$z = w \cdot x + b$$

$$\sigma(z) = \frac{1}{1 + exp(-z)}$$

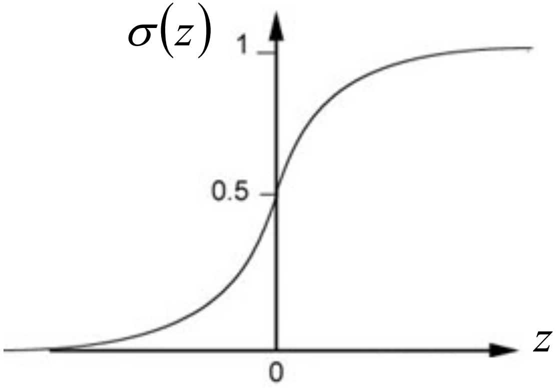

Function set:

$$f_{w,b}(x) = P_{w,b}(C_1|x)$$

Including all different w and b

Step 1: Function Set

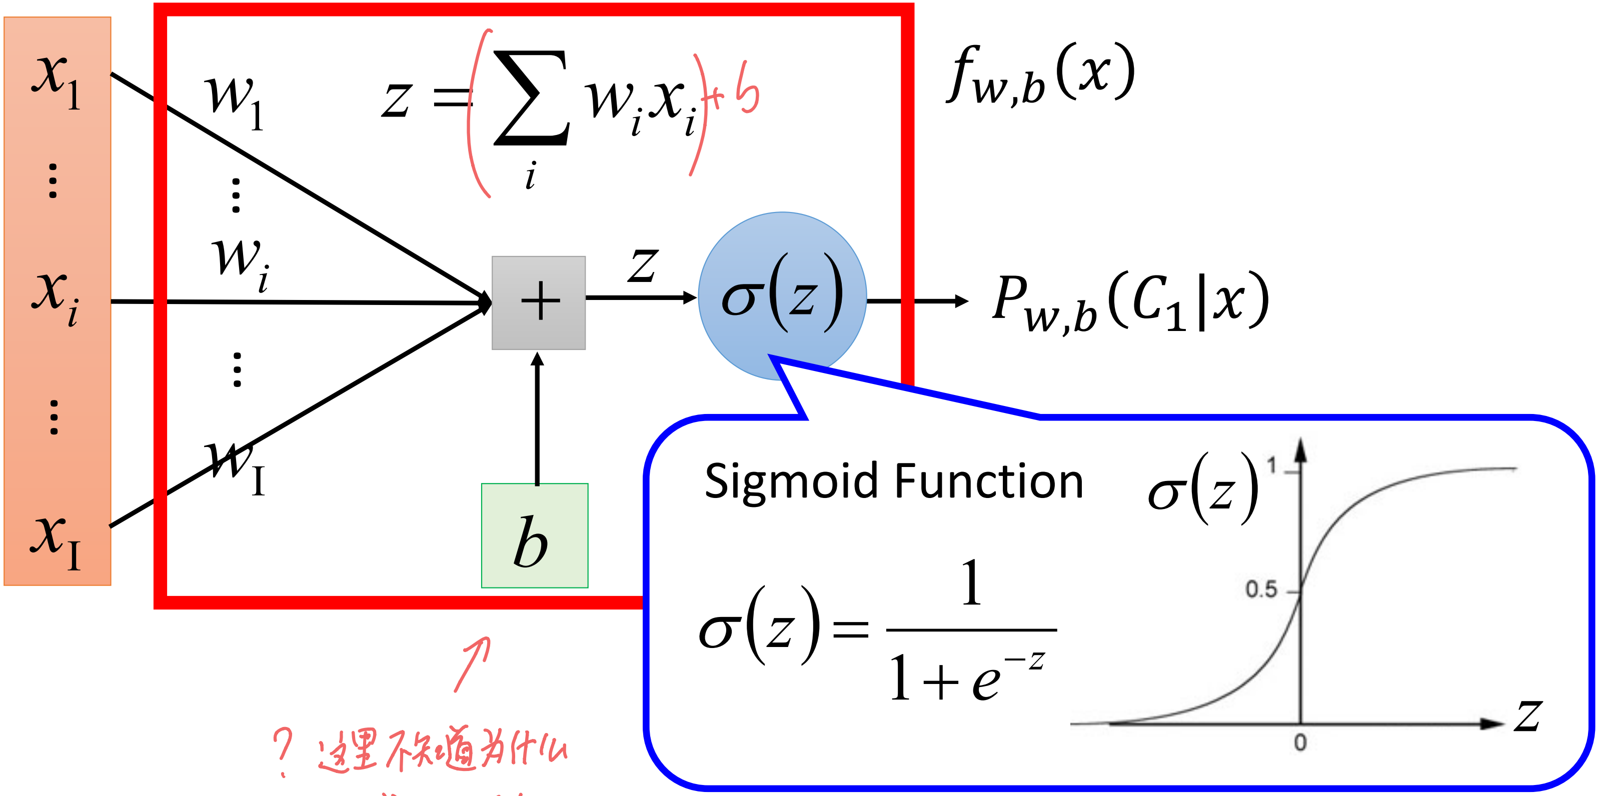

$$z = \left(\sum_i w_i x_i\right) + b$$

$$\sigma(z) = \frac{1}{1+e^{-z}}$$

? 这里不知道为什么
b变成1个维度的
(线性回归中, b是多个维度)

	Logistic Regression	***Linear Regression***
Step 1:	$f_{w,b}(x) = \underline{\underline{\sigma}}\left(\sum_i w_i x_i + b\right)$ Output: between 0 and 1	$f_{w,b}(x) = \sum_i w_i x_i + b$ Output: any value
Step 2:		
Step 3:		

Step 2: Goodness of a Function

Training Data

x^1	x^2	x^3	……	x^N
C_1	C_1	C_2		C_1

Assume the data is generated based on $f_{w,b}(x) = P_{w,b}(C_1|x)$

Given a set of w and b, what is its probability of generating the data?

$$L(w, b) = f_{w,b}(x^1) f_{w,b}(x^2) \left(1 - f_{w,b}(x^3)\right) \cdots f_{w,b}(x^N)$$

The most likely w* and b* is the one with the largest $L(w, b)$.

$$w^*, b^* = arg \max_{w,b} L(w, b)$$

x^1 x^2 x^3 ……

C_1 C_1 C_2

→

x^1 x^2 x^3 ……

$\hat{y}^1 = 1$ $\hat{y}^2 = \cancel{0}\,1$ $\hat{y}^3 = \cancel{1}\,0$

$\hat{y}^n$: 1 for class 1, 0 for class 2

$$L(w,b) = f_{w,b}(x^1) f_{w,b}(x^2)\left(1 - f_{w,b}(x^3)\right)\cdots$$

$$\boxed{w^*, b^* = arg\max_{w,b} L(w,b)} = \boxed{w^*, b^* = arg\min_{w,b} -lnL(w,b)}$$

这样好算

$$-lnL(w,b)$$

$$= -lnf_{w,b}(x^1) \rightarrow -[\,1\;lnf(x^1) \cancel{+\;0\;ln(1-f(x^1))}\,]$$

$$-lnf_{w,b}(x^2) \rightarrow -[\,1\;lnf(x^2) \cancel{+\;0\;ln(1-f(x^2))}\,]$$

$$-ln\left(1 - f_{w,b}(x^3)\right) \rightarrow -[\,\cancel{0\;lnf(x^3) +}\;1\;ln(1-f(x^3))\,]$$

$\vdots$

Step 2: Goodness of a Function

L= Likelihood

$$L(w,b) = f_{w,b}(x^1) f_{w,b}(x^2)\left(1 - f_{w,b}(x^3)\right) \cdots f_{w,b}(x^N)$$

$$-lnL(w,b) \overset{负号}{=} -\left(lnf_{w,b}(x^1) + lnf_{w,b}(x^2) + ln\left(1 - f_{w,b}(x^3)\right) \cdots\right)$$

$\hat{y}^n$: 1 for class 1, 0 for class 2

$$= \sum_n -\underbrace{\left[\overset{分类1}{\underline{\hat{y}^n}}\, \overset{分类1的可能性}{\underline{lnf_{w,b}(x^n)}} + \overset{分类2}{\underline{(1 - \hat{y}^n)}}\, ln\left(\overset{分类2的可能性 = 1- P(分类1)}{\underline{1 - f_{w,b}(x^n)}}\right)\right]}_{\text{Cross entropy between two Bernoulli distribution}}$$

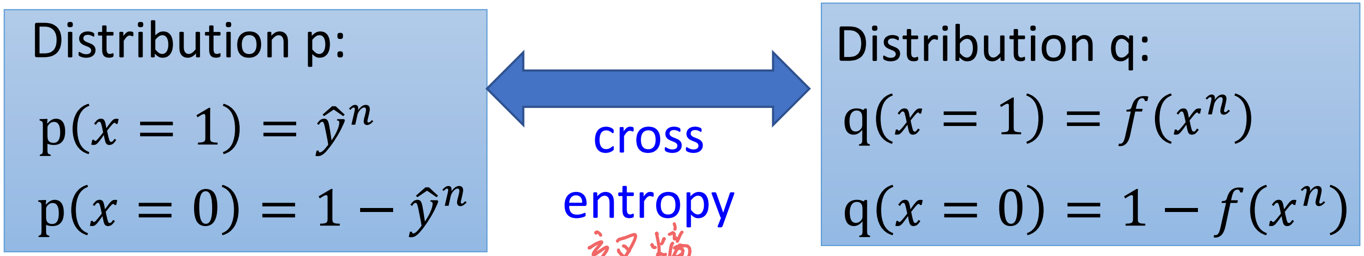

Distribution p:

$p(x = 1) = \hat{y}^n$

$p(x = 0) = 1 - \hat{y}^n$

cross entropy

交叉熵

Distribution q:

$q(x = 1) = f(x^n)$

$q(x = 0) = 1 - f(x^n)$

$$H(p,q) = -\sum_x p(x) ln\left(q(x)\right)$$

恒量 p.q 相似程度

Logistic Regression

Step 1: $$f_{w,b}(x) = \sigma\left(\sum_i w_i x_i + b\right)$$

Output: between 0 and 1

Step 2: Training data: $(x^n, \hat{y}^n)$

$\hat{y}^n$: 1 for class 1, 0 for class 2

$$L(f) = \sum_n C(f(x^n), \hat{y}^n)$$

Loss

Linear Regression

$$f_{w,b}(x) = \sum_i w_i x_i + b$$

Output: any value

Training data: $(x^n, \hat{y}^n)$

$\hat{y}^n$: a real number

$$L(f) = \frac{1}{2}\sum_n (f(x^n) - \hat{y}^n)^2$$

方便求导 (反正求最小，乘$\frac{1}{2}$无所谓)

Cross entropy:

$$C(f(x^n), \hat{y}^n) = -\left[\hat{y}^n ln f(x^n) + (1 - \hat{y}^n) ln\left(1 - f(x^n)\right)\right]$$

Question: Why don’t we simply use square error as linear regression?

下面用均方差试试 ↓

Step 3: Find the best function

$$\frac{-lnL(w,b)}{\partial w_i} = \sum_n -\left[\hat{y}^n \overset{\left(1-f_{w,b}(x^n)\right)x_i^n}{\boxed{\frac{lnf_{w,b}(x^n)}{\partial w_i}}} + (1-\hat{y}^n)\frac{ln\left(1-f_{w,b}(x^n)\right)}{\partial w_i}\right]$$

$$\Downarrow$$

$$\frac{\partial lnf_{w,b}(x)}{\partial w_i} = \frac{\partial lnf_{w,b}(x)}{\partial z}\frac{\partial z}{\partial w_i} \qquad \frac{\partial z}{\partial w_i} = x_i$$

链式法则

$$\frac{\partial ln\sigma(z)}{\partial z} = \frac{1}{\sigma(z)}\frac{\partial \sigma(z)}{\partial z} = \frac{1}{\cancel{\sigma(z)}}\cancel{\sigma(z)}\left(1-\sigma(z)\right)$$

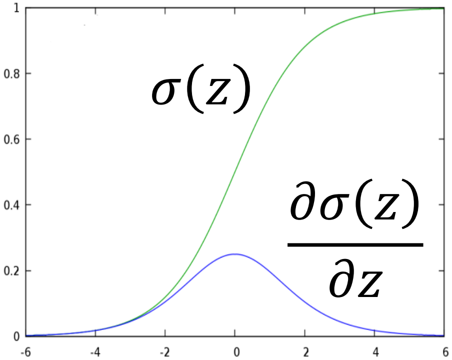

$$f_{w,b}(x) = \sigma(z)$$
$$= 1/1 + exp(-z)$$

$$z = w \cdot x + b = \sum_i w_i x_i + b$$

Step 3: Find the best function

$$\frac{-lnL(w,b)}{\partial w_i} = \sum_n -\left[\hat{y}^n \overset{\left(1-f_{w,b}(x^n)\right)x_i^n}{\boxed{\frac{lnf_{w,b}(x^n)}{\partial w_i}}} + (1-\hat{y}^n) \overset{f_{w,b}(x^n)x_i^n}{\boxed{\frac{ln\left(1-f_{w,b}(x^n)\right)}{\partial w_i}}}\right]$$

$$\frac{\partial ln\left(1-f_{w,b}(x)\right)}{\partial w_i} = \frac{\partial ln\left(1-f_{w,b}(x)\right)}{\partial z}\frac{\partial z}{\partial w_i} \qquad \frac{\partial z}{\partial w_i} = x_i$$

$$\frac{\partial ln\left(1-\sigma(z)\right)}{\partial z} = -\frac{1}{1-\sigma(z)}\frac{\partial \sigma(z)}{\partial z} = -\frac{1}{\cancel{1-\sigma(z)}}\sigma(z)\cancel{\left(1-\sigma(z)\right)}$$

$f_{w,b}(x) = \sigma(z)$

$= 1/1+exp(-z)$

$$z = w \cdot x + b = \sum_i w_i x_i + b$$

Step 3: Find the best function

$$\frac{-lnL(w,b)}{\partial w_i} = \sum_n -\left[\hat{y}^n \overset{\left(1-f_{w,b}(x^n)\right)x_i^n}{\boxed{\frac{lnf_{w,b}(x^n)}{\partial w_i}}} + (1-\hat{y}^n)\overset{f_{w,b}(x^n)x_i^n}{\boxed{\frac{ln\left(1-f_{w,b}(x^n)\right)}{\partial w_i}}}\right]$$

$$= \sum_n -\left[\hat{y}^n \underline{\left(1-f_{w,b}(x^n)\right)\underline{x_i^n}} - (1-\hat{y}^n)\underline{f_{w,b}(x^n)\underline{x_i^n}}\right]$$

$$= \sum_n -\left[\hat{y}^n \cancel{-\hat{y}^n f_{w,b}(x^n)} - f_{w,b}(x^n) \cancel{+\hat{y}^n f_{w,b}(x^n)}\right] \underline{x_i^n}$$

$$= \sum_n -\left(\underline{\hat{y}^n - f_{w,b}(x^n)}\right)x_i^n \leftarrow \text{梯度}$$

Larger difference,
larger update

$$w_i \leftarrow w_i - \underset{= \text{学习率}}{\eta} \left[\sum_n -\left(\hat{y}^n - f_{w,b}(x^n)\right)x_i^n\right]$$

	Logistic Regression	***Linear Regression***
Step 1:	$f_{w,b}(x) = \sigma\left(\sum_i w_i x_i + b\right)$ Output: between 0 and 1	$f_{w,b}(x) = \sum_i w_i x_i + b$ Output: any value
Step 2:	Training data: $(x^n, \hat{y}^n)$ $\hat{y}^n$: 1 for class 1, 0 for class 2 $L(f) = \sum_n C(f(x^n), \hat{y}^n)$	Training data: $(x^n, \hat{y}^n)$ $\hat{y}^n$: a real number $L(f) = \frac{1}{2}\sum_n (f(x^n) - \hat{y}^n)^2$

Step 3:

Logistic regression: $w_i \leftarrow w_i - \eta \sum_n -\underline{\left(\hat{y}^n - f_{w,b}(x^n)\right)} x_i^n$

Linear regression: $w_i \leftarrow w_i - \eta \sum_n -\underline{\left(\hat{y}^n - f_{w,b}(x^n)\right)} x_i^n$

Logistic Regression + Square Error

Step 1: $$f_{w,b}(x) = \sigma\left(\sum_i w_i x_i + b\right)$$

Step 2: Training data: $(x^n, \hat{y}^n)$, $\hat{y}^n$: 1 for class 1, 0 for class 2

$$L(f) = \frac{1}{2}\sum_n \left(f_{w,b}(x^n) - \hat{y}^n\right)^2$$

Step 3:

$$\frac{\partial\,(f_{w,b}(x) - \hat{y})^2}{\partial w_i} = 2\left(f_{w,b}(x) - \hat{y}\right)\frac{\partial f_{w,b}(x)}{\partial z}\frac{\partial z}{\partial w_i}$$

$$= 2\left(f_{w,b}(x) - \hat{y}\right) f_{w,b}(x)\left(1 - f_{w,b}(x)\right) x_i$$

$\hat{y}^n = 1$

If $f_{w,b}(x^n) = 1$ (close to target) → $\partial L / \partial w_i = 0$

If $f_{w,b}(x^n) = 0$ (far from target) → $\partial L / \partial w_i = 0$

Logistic Regression + Square Error

Step 1: $f_{w,b}(x) = \sigma\left(\sum_i w_i x_i + b\right)$

Step 2: Training data: $(x^n, \hat{y}^n)$, $\hat{y}^n$: 1 for class 1, 0 for class 2

$$L(f) = \frac{1}{2}\sum_n \left(f_{w,b}(x^n) - \hat{y}^n\right)^2$$

Step 3:

$$\frac{\partial\,(f_{w,b}(x) - \hat{y})^2}{\partial w_i} = 2\left(f_{w,b}(x) - \hat{y}\right)\frac{\partial f_{w,b}(x)}{\partial z}\frac{\partial z}{\partial w_i}$$

$$= 2\left(f_{w,b}(x) - \hat{y}\right) f_{w,b}(x)\left(1 - f_{w,b}(x)\right) x_i$$

$\hat{y}^n = 0$ If $f_{w,b}(x^n) = 1$ (far from target) ➡ $\partial L / \partial w_i = 0$

If $f_{w,b}(x^n) = 0$ (close to target) ➡ $\partial L / \partial w_i = 0$

因为 y 的 output 不是 0 就是 1，所以用均方差时，会导致梯度消失

均方差 .

Cross Entropy v.s. Square Error

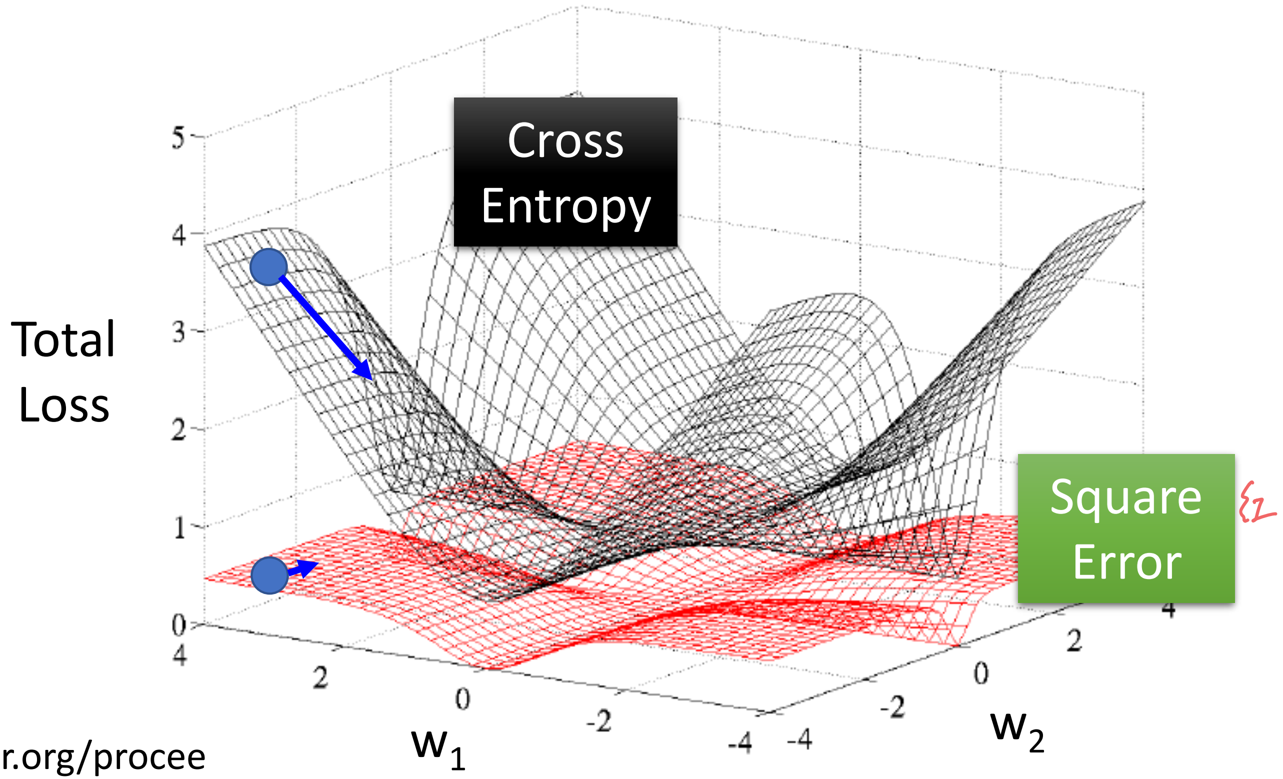

http://jmlr.org/proceedings/papers/v9/glorot10a/glorot10a.pdf

Discriminative v.s. Generative

判别模型

生成模型

$$P(C_1|x) = \sigma(w \cdot x + b)$$

directly find ***w*** and b

Find $\mu^1, \mu^2, \Sigma^{-1}$

$$w^T = (\mu^1 - \mu^2)^T \Sigma^{-1}$$

$$b = -\frac{1}{2}(\mu^1)^T(\Sigma^1)^{-1}\mu^1 + \frac{1}{2}(\mu^2)^T(\Sigma^2)^{-1}\mu^2 + ln\frac{N_1}{N_2}$$

Will we obtain the same set of w and b?

← 假设不同 →

无假设

假设~高斯分布

The same model (function set), but different function is selected by the same training data.

Generative v.s. Discriminative

Gaussian

Generative

Discriminative

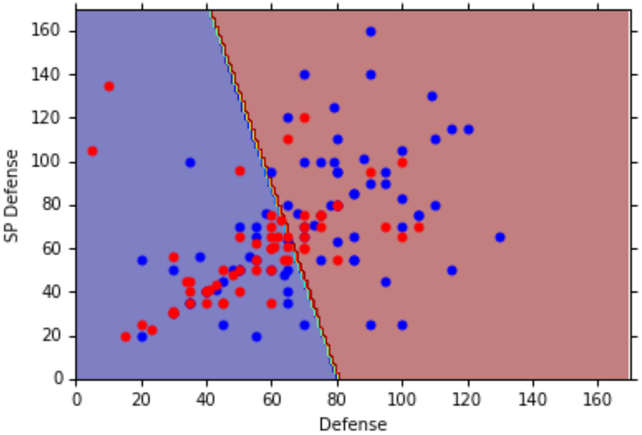

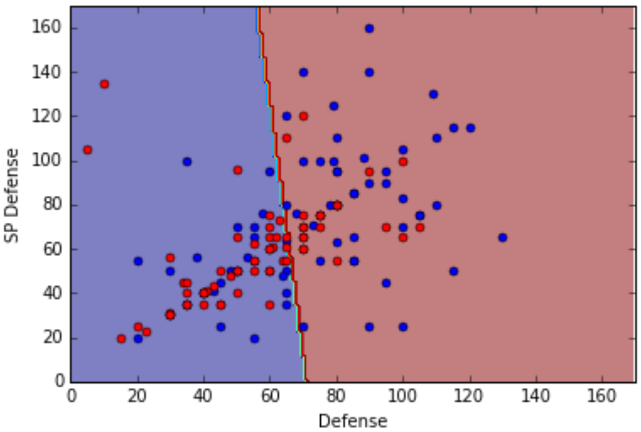

All: total, hp, att, sp att, de, sp de, speed

73% accuracy

79% accuracy

Generative v.s. Discriminative

- Example

Training Data

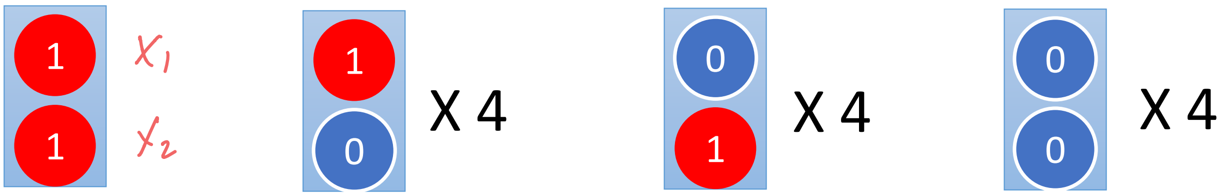

Class 1 Class 2 Class 2 Class 2

Testing Data

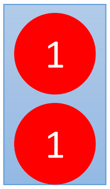

Class 1?
Class 2?

朴素贝叶斯

How about Naïve Bayes?

$$P(x|C_i) = P(x_1|C_i)P(x_2|C_i)$$

Generative v.s. Discriminative

- Example

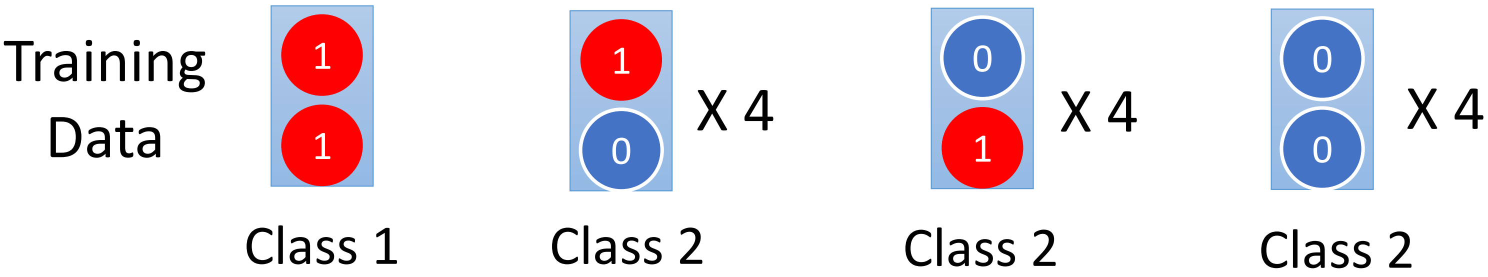

Training Data

Class 1

Class 2 X 4

Class 2 X 4

Class 2 X 4

$$P(C_1) = \frac{1}{13} \qquad P(x_1 = 1|C_1) = 1 \qquad P(x_2 = 1|C_1) = 1$$

$$P(C_2) = \frac{12}{13} \qquad P(x_1 = 1|C_2) = \frac{1}{3} \qquad P(x_2 = 1|C_2) = \frac{1}{3}$$

Training Data

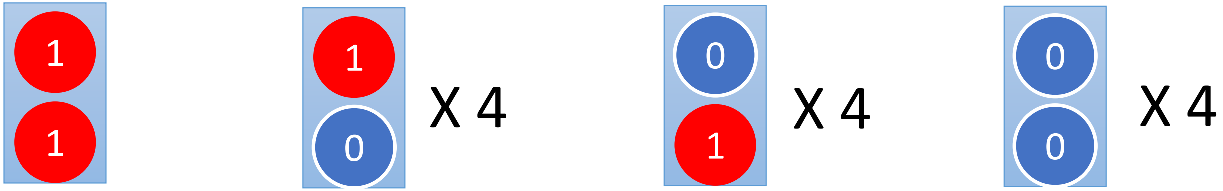

Class 1 Class 2 Class 2 Class 2

样本少，偏离μ周围

概率分布，给了数据每个样本一个权重

Testing Data

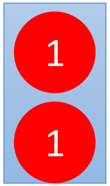

$$\overset{<0.5}{P(C_1|x)} = \frac{\overbrace{P(x|C_1)}^{1\times 1}\overbrace{P(C_1)}^{\frac{1}{13}}}{\underbrace{P(x|C_1)}_{1\times 1}\underbrace{P(C_1)}_{\frac{1}{13}} + \underbrace{P(x|C_2)}_{\frac{1}{3}\times\frac{1}{3}}\underbrace{P(C_2)}_{\frac{12}{13}}}$$

$$P(C_1) = \frac{1}{13} \qquad P(x_1 = 1|C_1) = 1 \qquad P(x_2 = 1|C_1) = 1$$

$$P(C_2) = \frac{12}{13} \qquad P(x_1 = 1|C_2) = \frac{1}{3} \qquad P(x_2 = 1|C_2) = \frac{1}{3}$$

Generative v.s. Discriminative

- Benefit of generative model
  - With the assumption of probability distribution, less training data is needed
  - With the assumption of probability distribution, more robust to the noise
  - Priors and class-dependent probabilities can be estimated from different sources.

[Bishop, P209-210]

Multi-class Classification (3 classes as example)

C$_1$: w^1, b_1 $\quad z_1 = w^1 \cdot x + b_1$

C$_2$: w^2, b_2 $\quad z_2 = w^2 \cdot x + b_2$

C$_3$: w^3, b_3 $\quad z_3 = w^3 \cdot x + b_3$

Probability:

- ■ $1 > y_i > 0$
- ■ $\sum_i y_i = 1$

$$y_i = P(C_i \mid x)$$

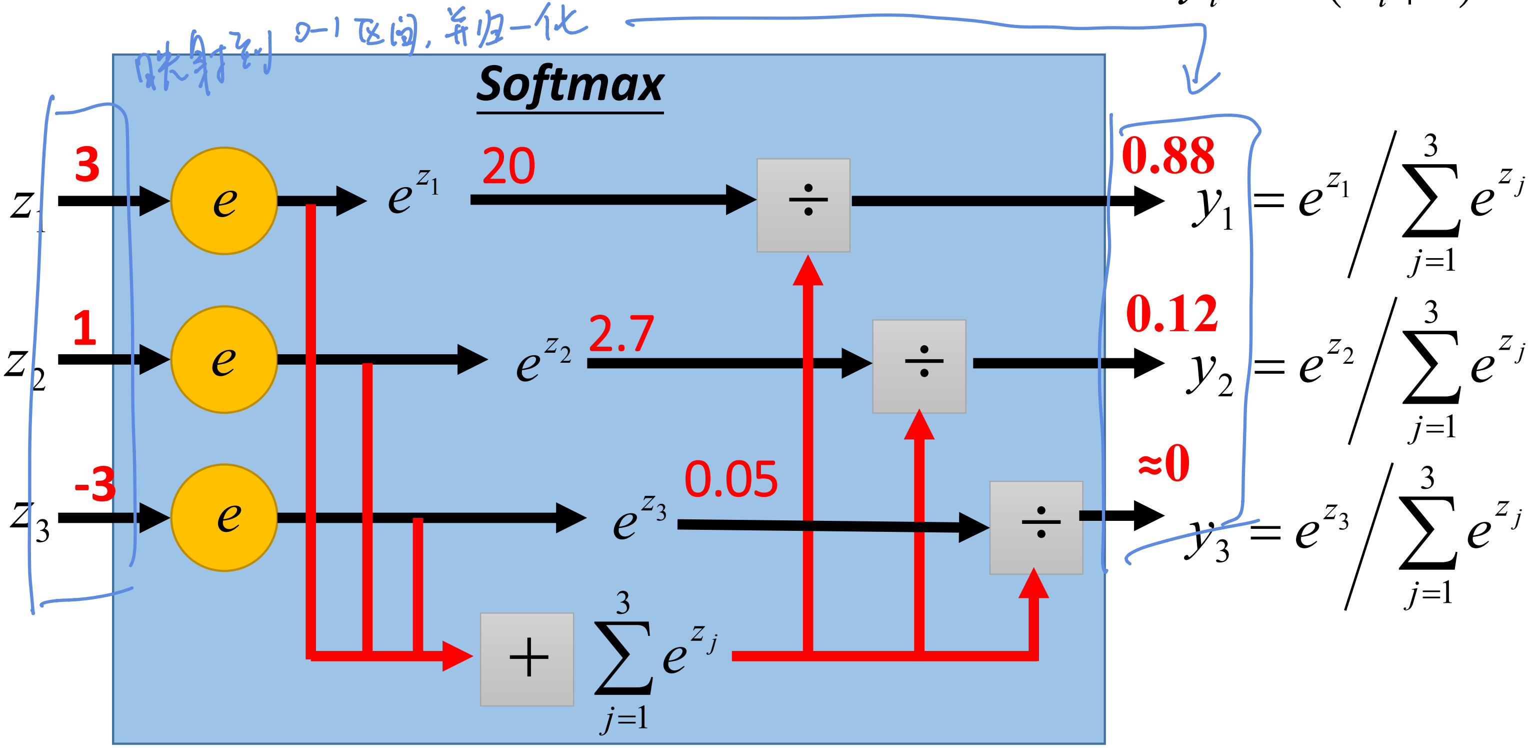

Multi-class Classification (3 classes as example)

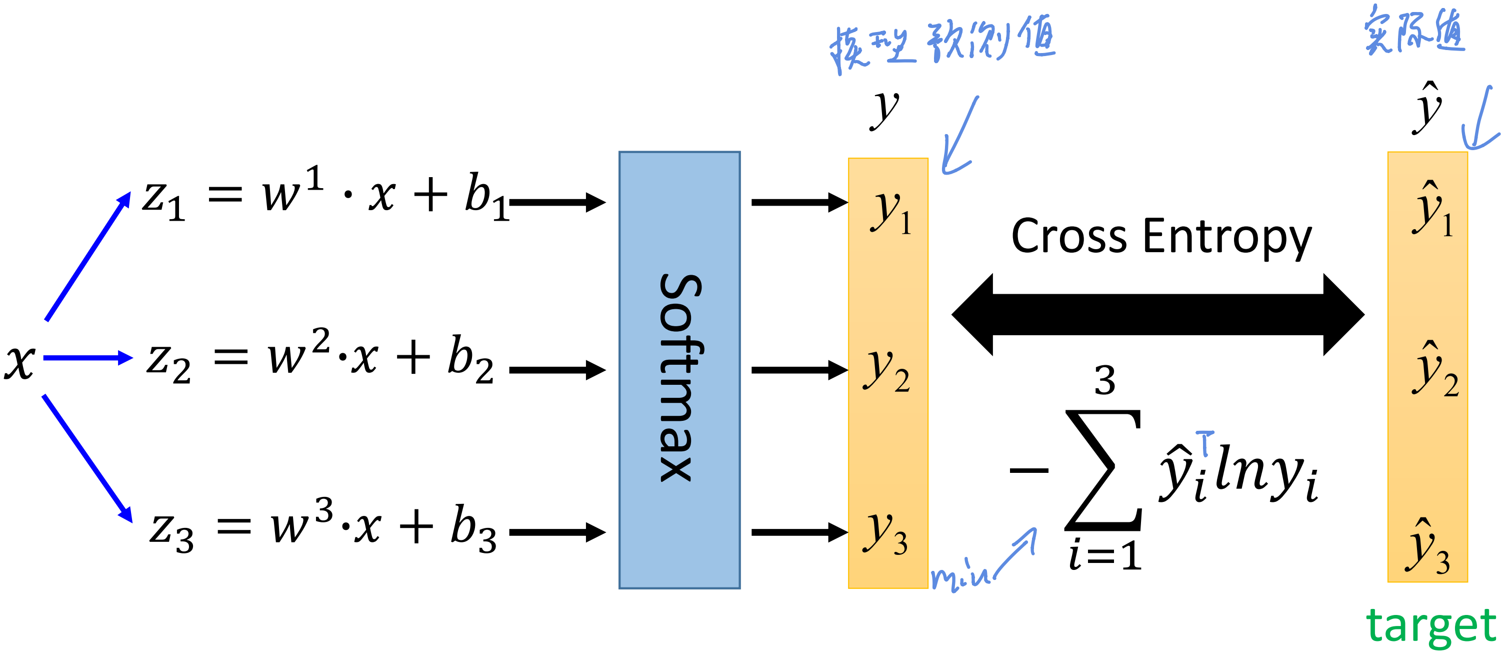

模型预测值

实际值

y

$\hat{y}$

x

$z_1 = w^1 \cdot x + b_1$

$z_2 = w^2 \cdot x + b_2$

$z_3 = w^3 \cdot x + b_3$

Softmax

y_1

y_2

y_3

Cross Entropy

$$-\sum_{i=1}^{3} \hat{y}_i^{T} \ln y_i$$

min

$\hat{y}_1$

$\hat{y}_2$

$\hat{y}_3$

target

If x ∈ class 1	If x ∈ class 2	If x ∈ class 3
$\hat{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$	$\hat{y} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$	$\hat{y} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$

Limitation of Logistic Regression

$$z = w_1 x_1 + w_2 x_2 + b$$

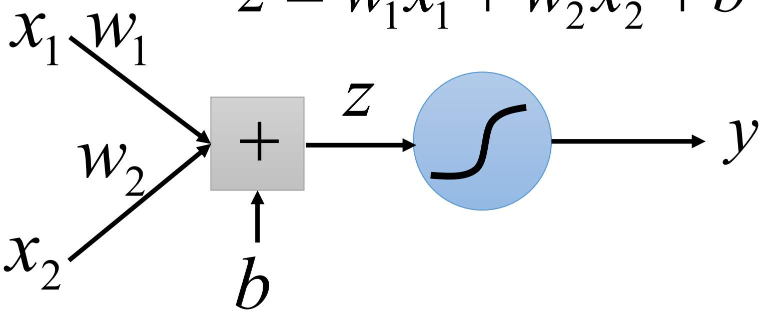

$$\begin{cases} Class1 & y \geq 0.5 \\ Class2 & y < 0.5 \end{cases}$$

Input Feature		Label
x_1	x_2	
0	0	Class 2
0	1	Class 1
1	0	Class 1
1	1	Class 2

Can we?

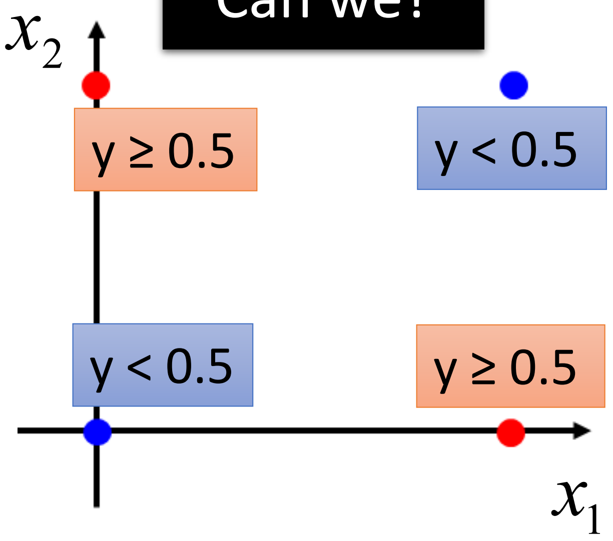

Limitation of Logistic Regression

- No, we can't ……

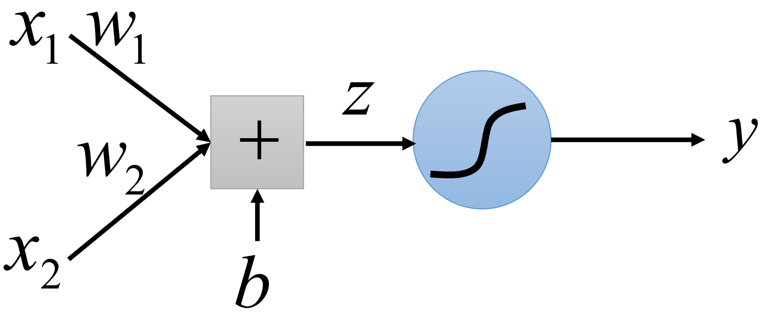

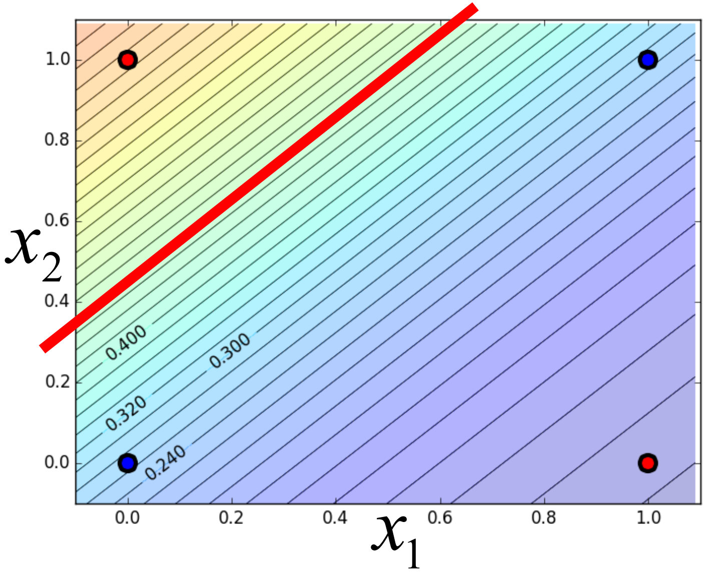

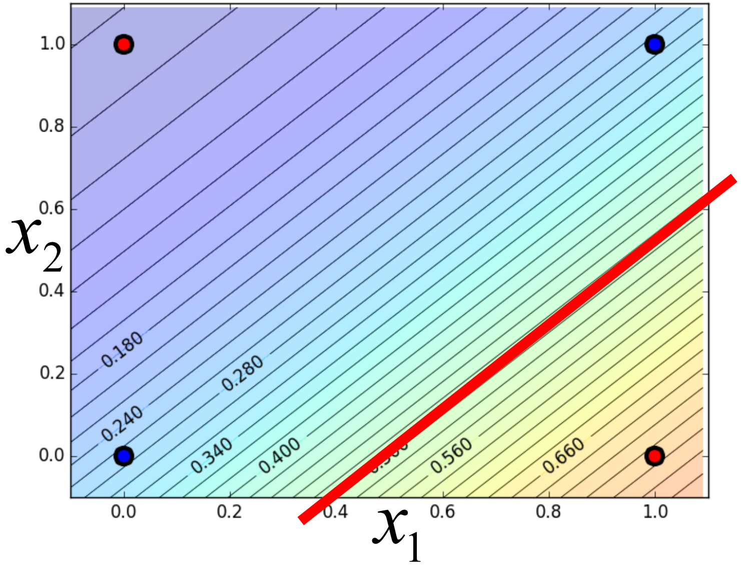

Limitation of Logistic Regression

- Feature Transformation

x_1': distance to $\begin{bmatrix}0\\0\end{bmatrix}$

x_2': distance to $\begin{bmatrix}1\\1\end{bmatrix}$

$$\begin{bmatrix}x_1\\x_2\end{bmatrix} \longrightarrow \begin{bmatrix}x_1'\\x_2'\end{bmatrix}$$

Not always easy to find a good transformation

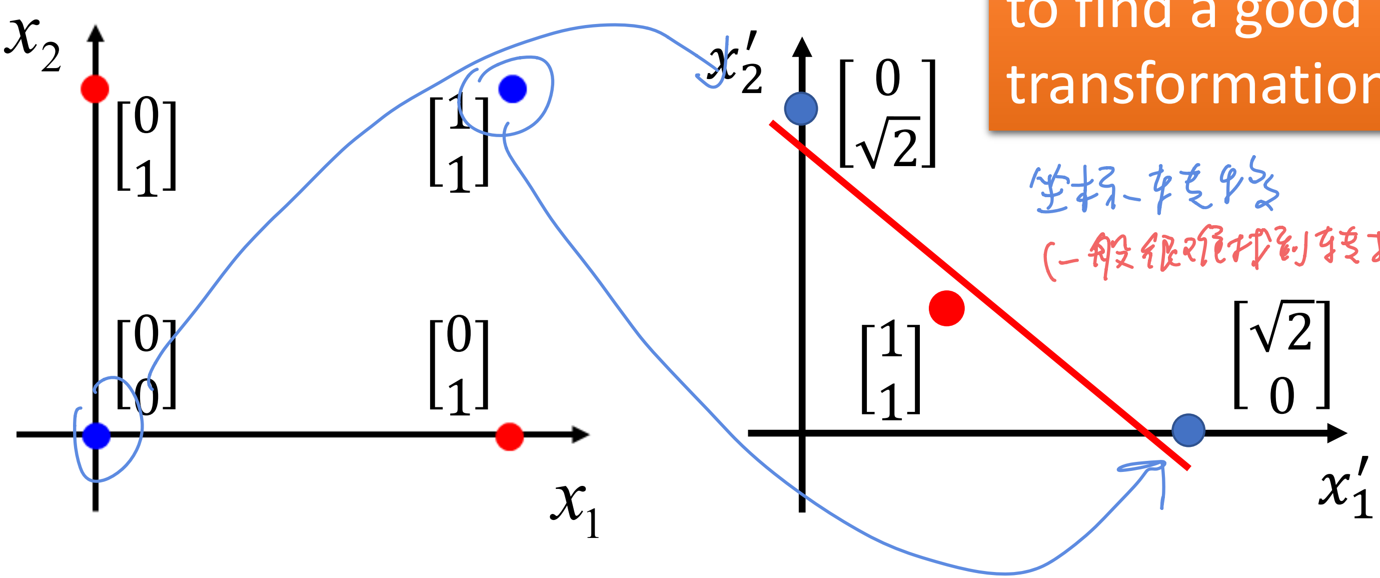

x_2
$\begin{bmatrix}0\\1\end{bmatrix}$
$\begin{bmatrix}1\\1\end{bmatrix}$
$\begin{bmatrix}0\\0\end{bmatrix}$
$\begin{bmatrix}0\\1\end{bmatrix}$
x_1
x_2'
$\begin{bmatrix}0\\\sqrt{2}\end{bmatrix}$
$\begin{bmatrix}1\\1\end{bmatrix}$
$\begin{bmatrix}\sqrt{2}\\0\end{bmatrix}$
x_1'

坐标转换

(一般很难找到转换规则)

Limitation of Logistic Regression

- Cascading logistic regression models

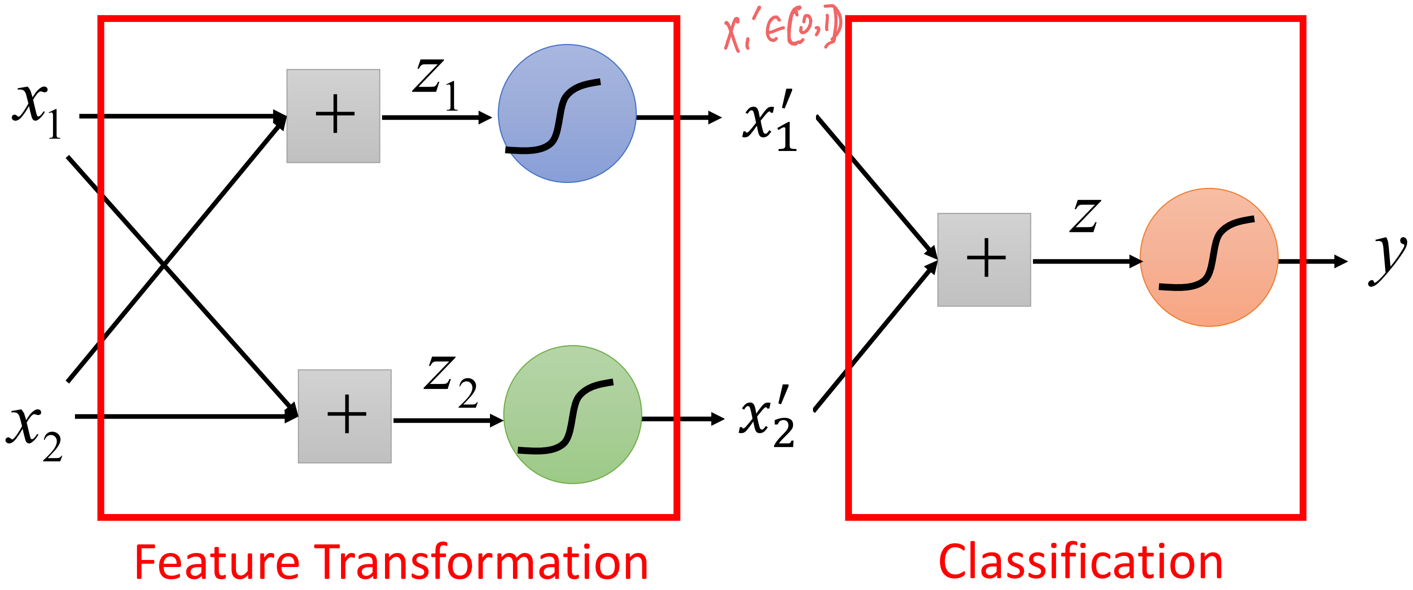

(ignore bias in this figure)

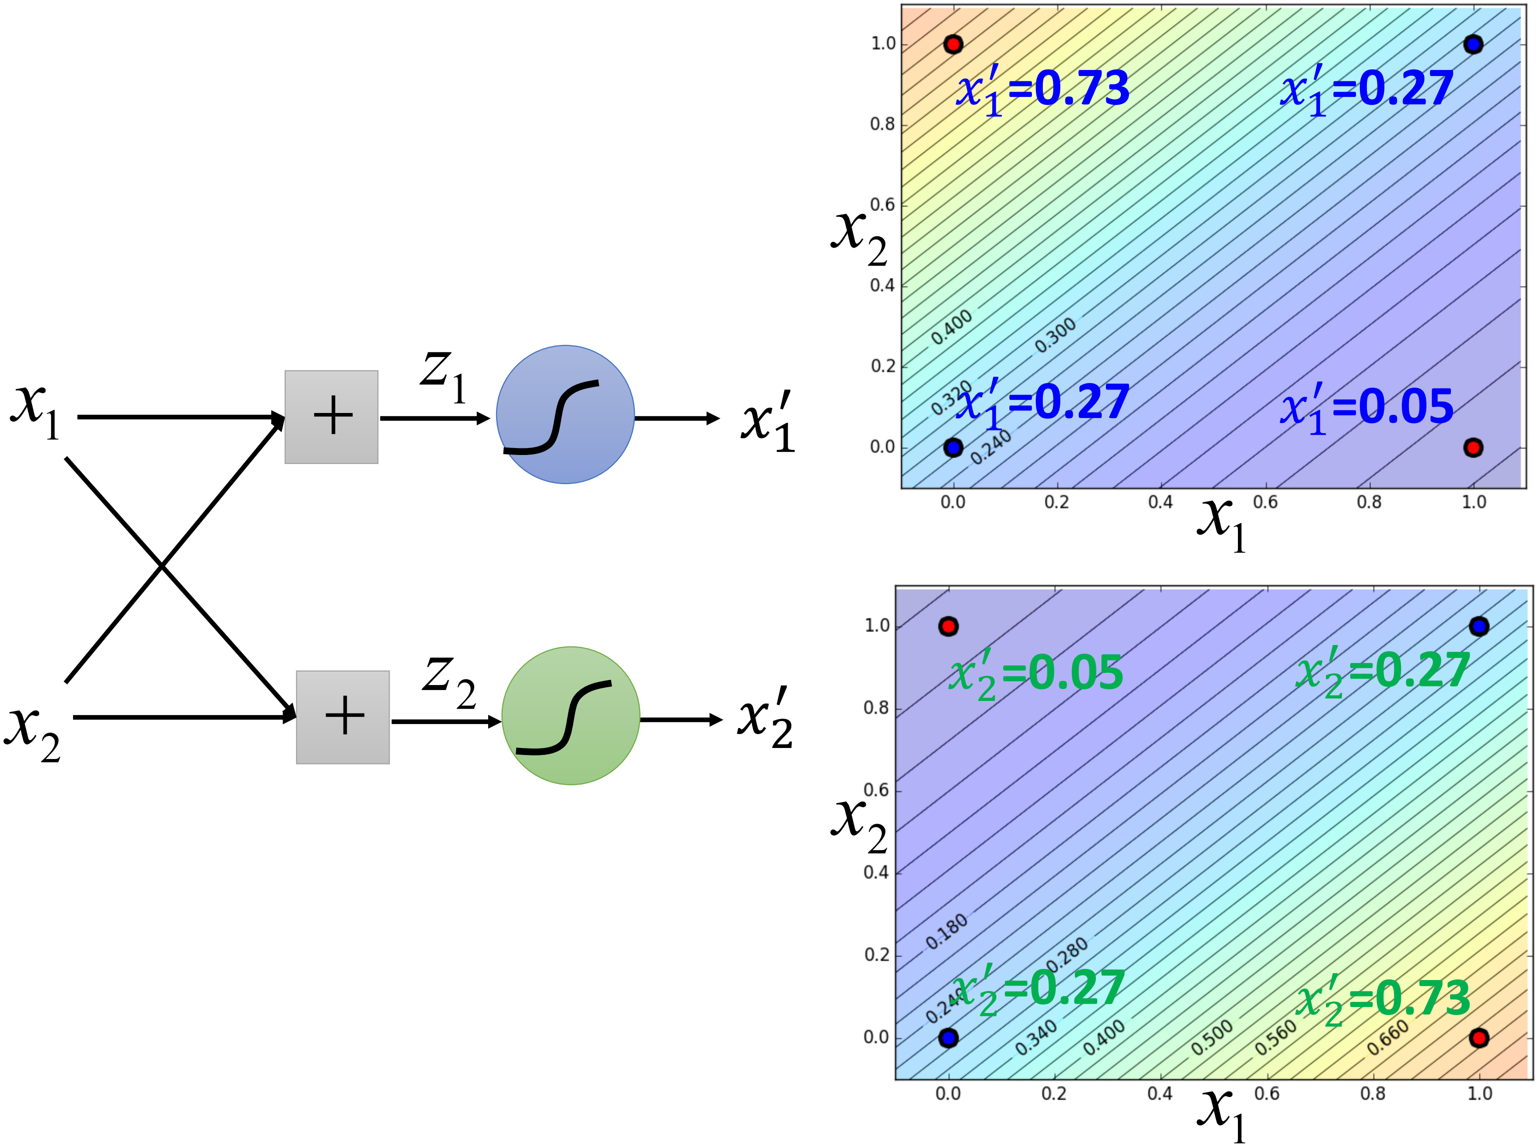
x_1
x_2
+
+
z_1
z_2
x'_1
x'_2
x'_1=0.73
x'_1=0.27
x'_1=0.27
x'_1=0.05
0.400
0.300
0.320
0.240
1.0
0.8
0.6
0.4
0.2
0.0
x_2
0.0
0.2
0.4
0.6
0.8
1.0
x_1
x'_2=0.05
x'_2=0.27
x'_2=0.27
x'_2=0.73
0.180
0.280
0.240
0.340
0.400
0.500
0.560
0.660
1.0
0.8
0.6
0.4
0.2
0.0
x_2
0.0
0.2
0.4
0.6
0.8
1.0
x_1

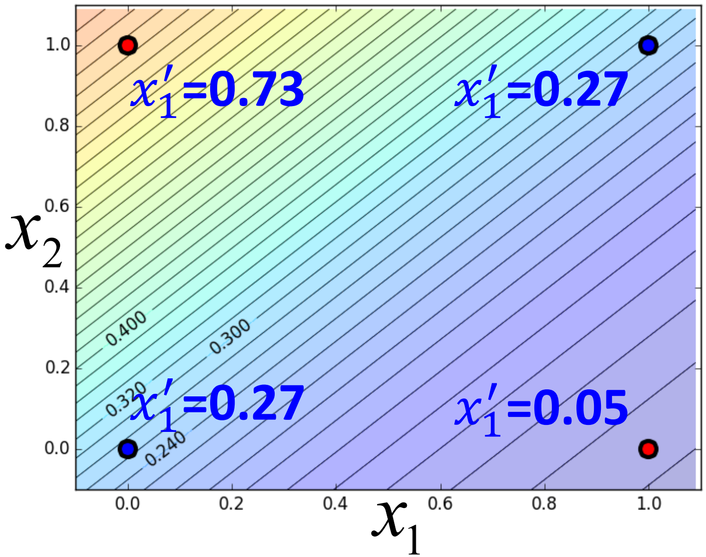
1.0
0.8
0.6
0.4
0.2
0.0
x_2
x'_1=0.73
x'_1=0.27
0.400
0.300
0.320
0.240
x'_1=0.27
x'_1=0.05
0.0
0.2
0.4
0.6
0.8
1.0
x_1

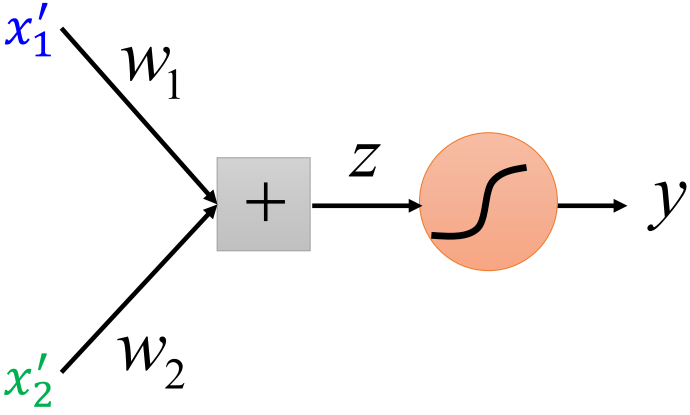
x'_1
w_1
+
z
y
x'_2
w_2

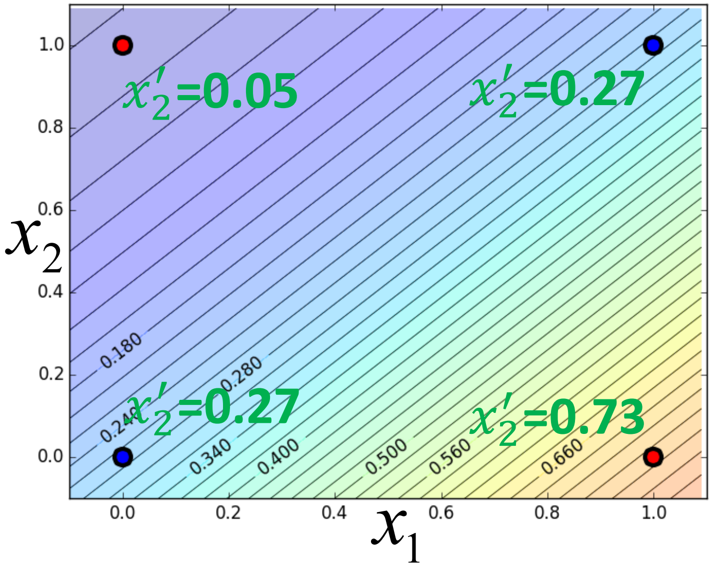
1.0
0.8
0.6
0.4
0.2
0.0
x_2
x'_2=0.05
x'_2=0.27
0.180
0.280
0.240
x'_2=0.27
x'_2=0.73
0.340
0.400
0.500
0.560
0.660
0.0
0.2
0.4
0.6
0.8
1.0
x_1

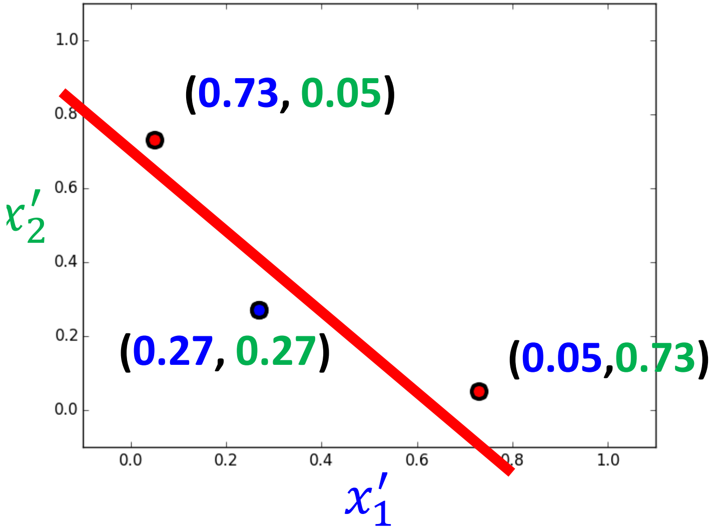
1.0
0.8
0.6
0.4
0.2
0.0
x'_2
(0.73, 0.05)
(0.27, 0.27)
(0.05,0.73)
0.0
0.2
0.4
0.6
0.8
1.0
x'_1

Deep Learning!

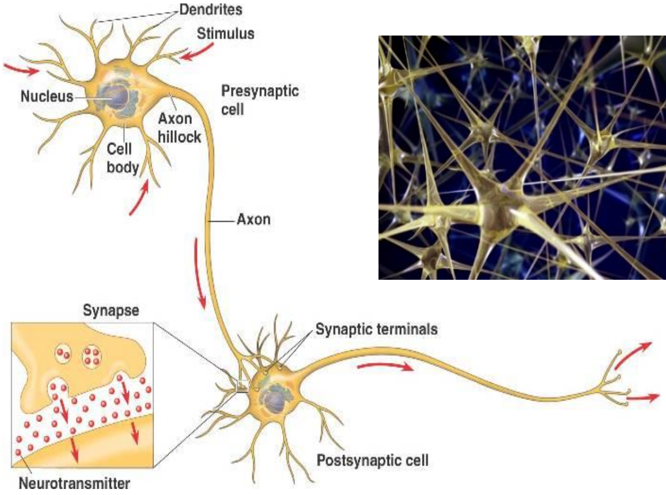

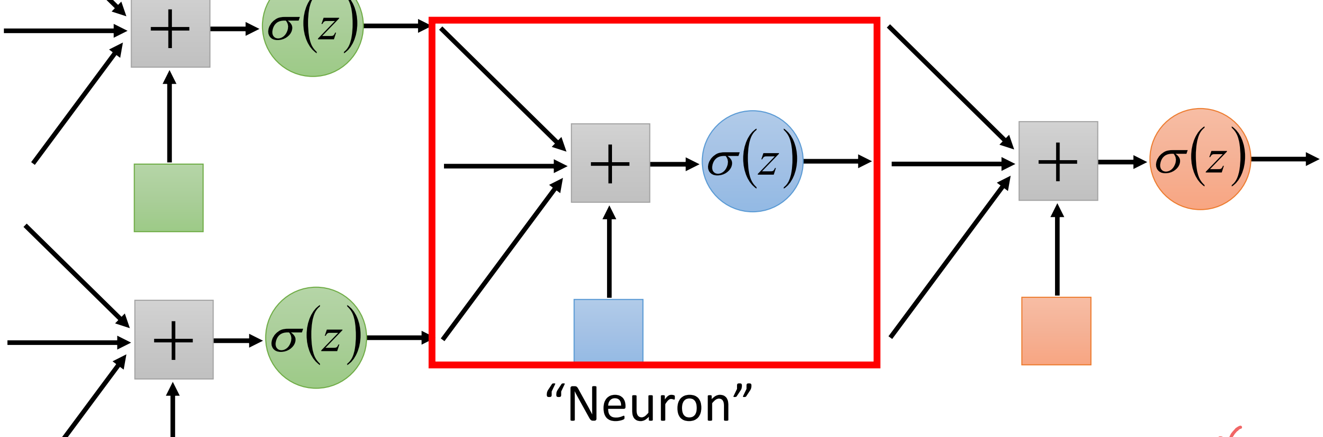

Neural Network

神经网络

Reference

- Bishop: Chapter 4.3